

Vectorizing Divergent Control-Flow for SIMD Applications


Fabian Wahlster

Technische Universität München

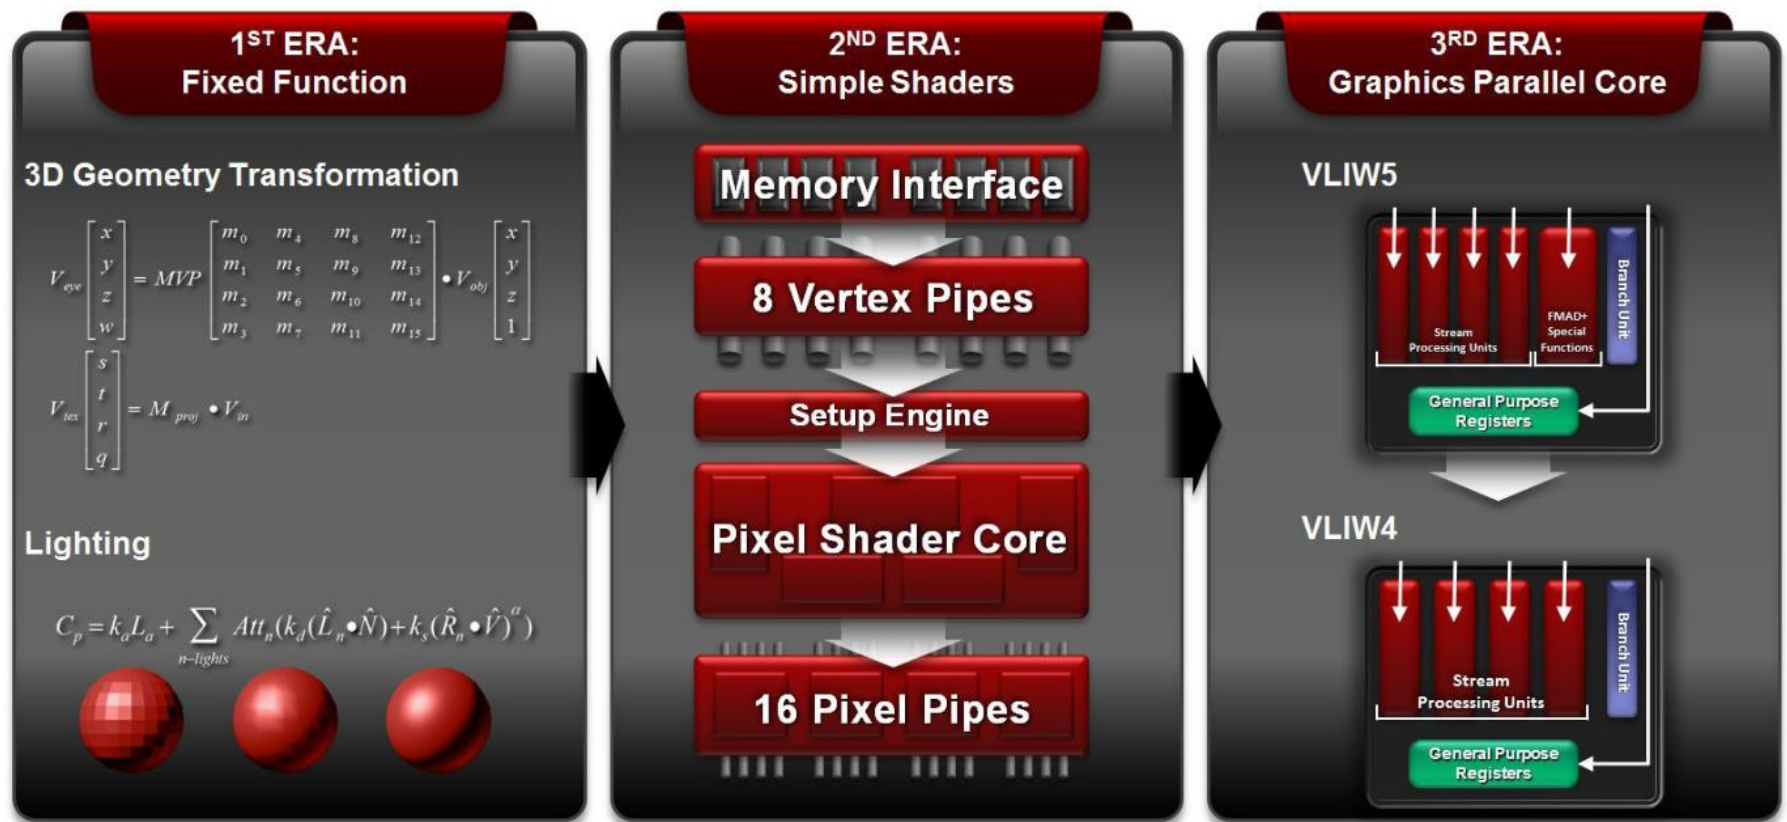
Fakultät für Informatik

Lehrstuhl für Rechnerarchitektur & Parallele Systeme

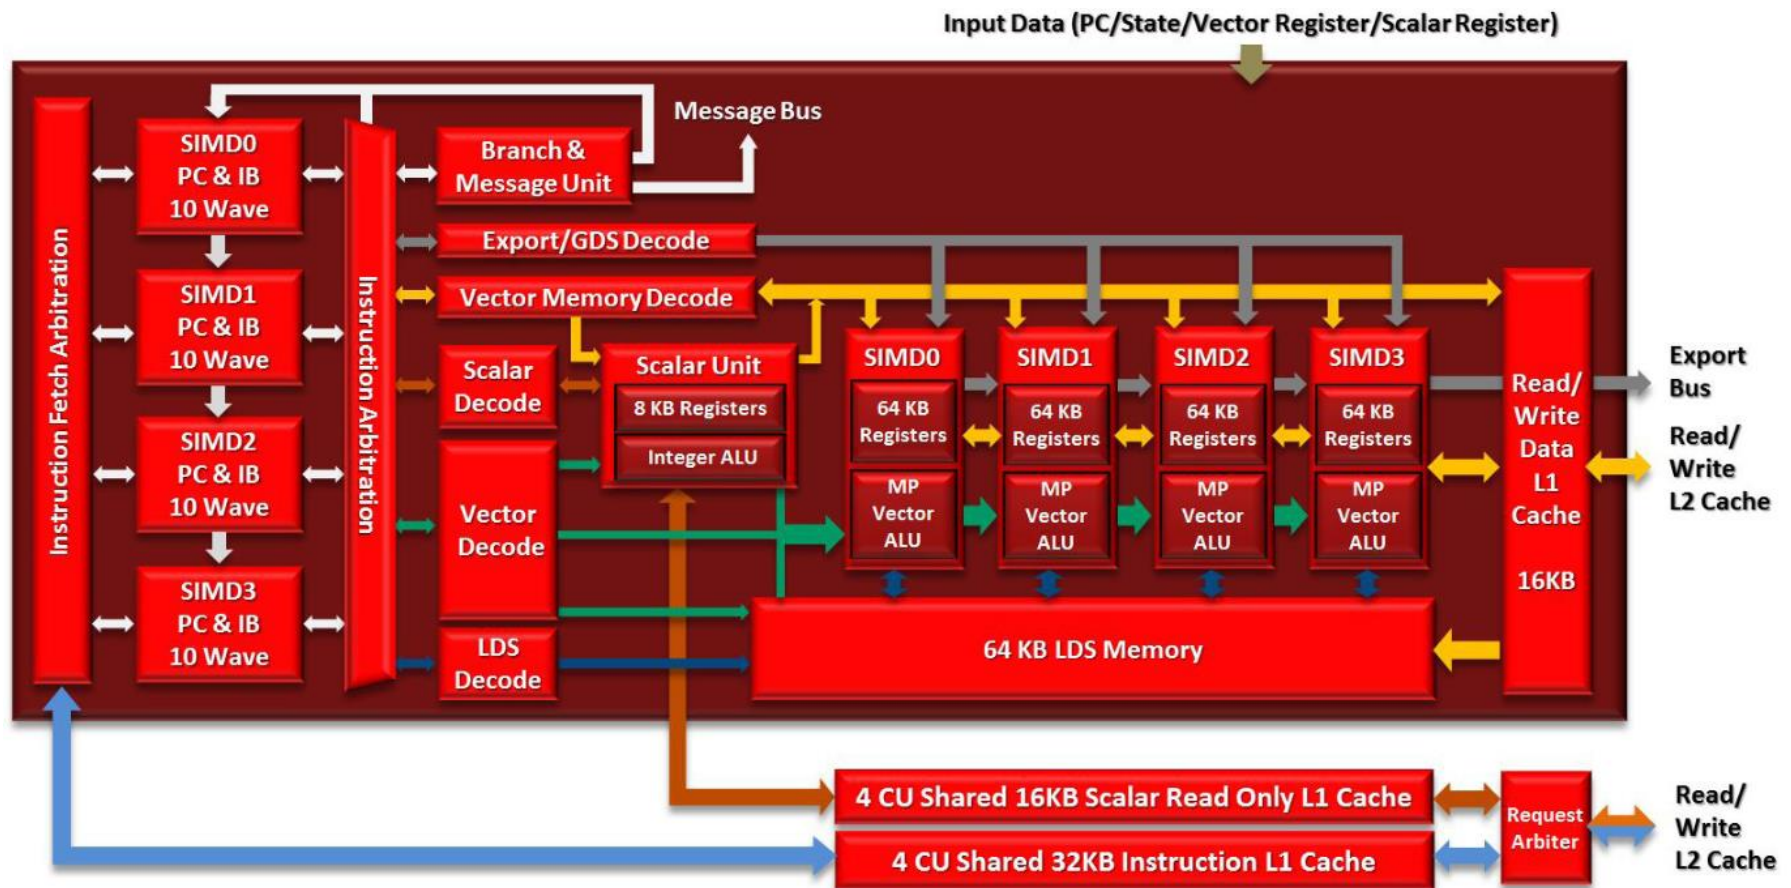
Garching, 14. Februar 2019



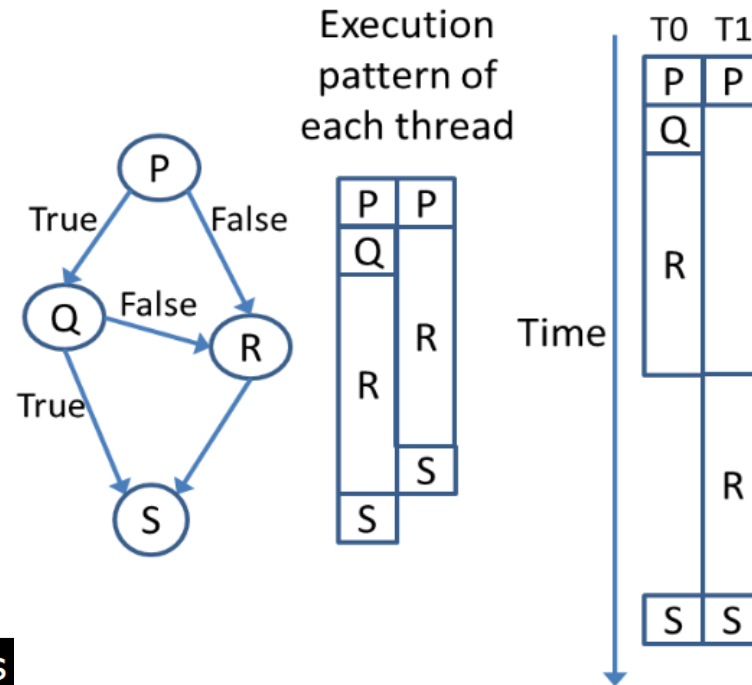
Hardware: Fixed function graphics to GPGPU



Hardware: Fixed function graphics to GPGPU



Problem Description: Divergence on wide SIMD



All lanes / invocations active

Subset of lanes / invocations active


wave / subgroup

if() statement

only triggers for some lanes

wave / subgroup

Problem Description: Converting thread-level code to wave-level ISA



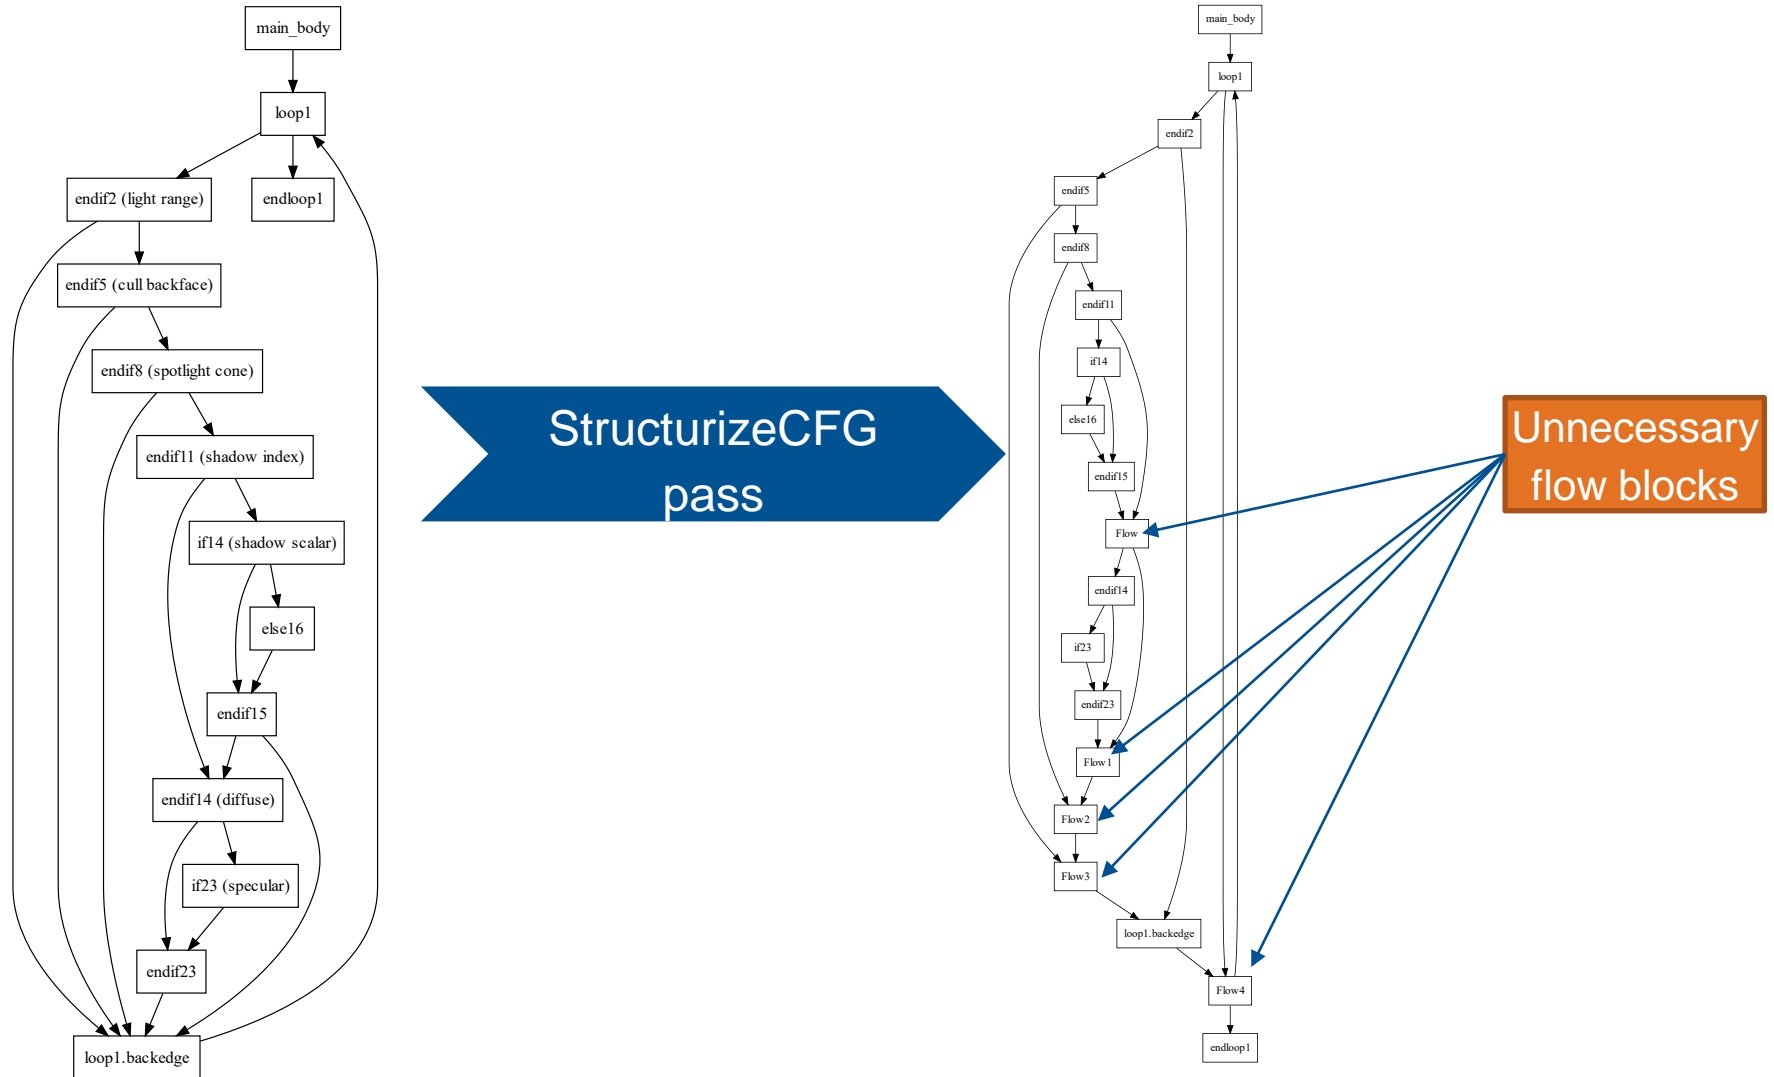
```
float fn0(float a, float b)
{
    if(a>b)
        return((a-b)*a);
    else
        return((b-a)*b);
}
```

```
//Registers r0 contains "a", r1 contains "b"
//Value is returned in r2
```

```
v_cmp_gt_f32    r0,r1        //a > b, establish VCC
s_mov_b64       s0,exec      //Save current exec mask
s_and_b64       exec,vcc,exec //Do "if"
s_cbranch_vccz  label0       //Branch if all lanes fail
v_sub_f32       r2,r0,r1     //result = a - b
v_mul_f32       r2,r2,r0     //result=result * a
```

```
label0:
s_andn2_b64     exec,s0,exec  //Do "else" (s0 & !exec)
s_cbranch_execz label1       //Branch if all lanes fail
v_sub_f32       r2,r1,r0     //result = b - a
v_mul_f32       r2,r2,r1     //result = result * b
label1:
s_mov_b64       exec,s0      //Restore exec mask
```

Problem Description: Structurization in LLVM



Related Works

Hardware Reconvergence (Stacks):



- Thread frontiers (Intel)
- Simultaneous Branch/Warp Interweaving (Nvidia)

Software:

- Partial Control flow linearization
- Taming Control Divergence in GPUs through Control Flow Linearization

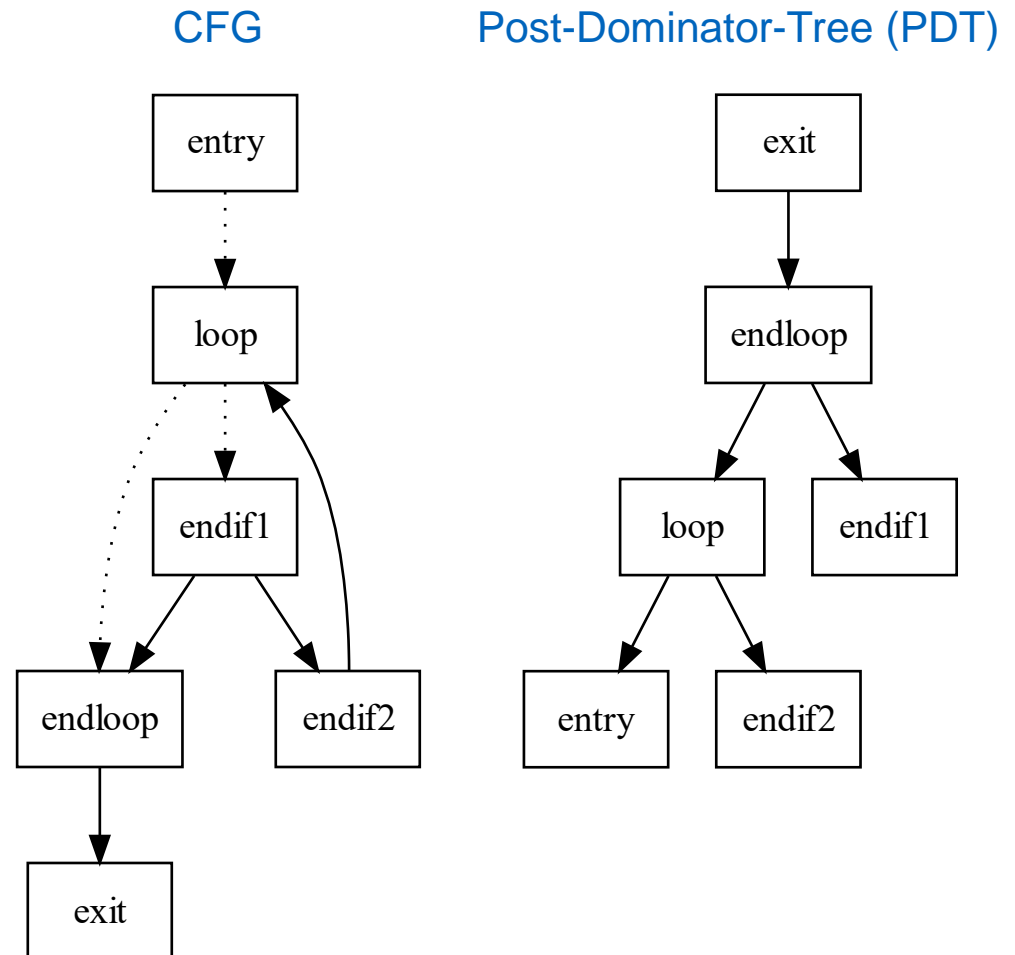
Dominator Trees

- Node **P** dominates node **Q** if every path from the **entry** node has to pass through **P** to connect to **Q**.



Post Dominator Trees

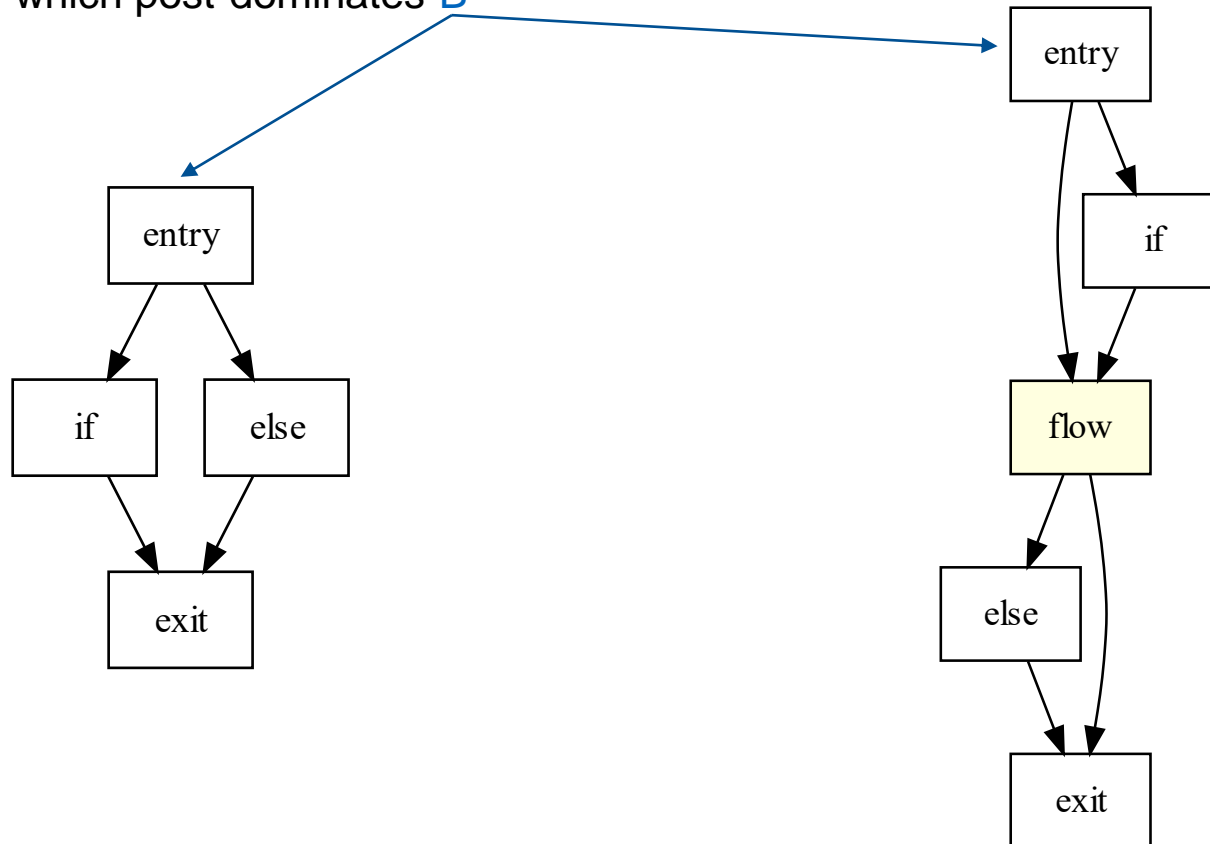
- Node **P** post-dominates node **Q** if every path from **Q** to the **exit** node has to through **P**.



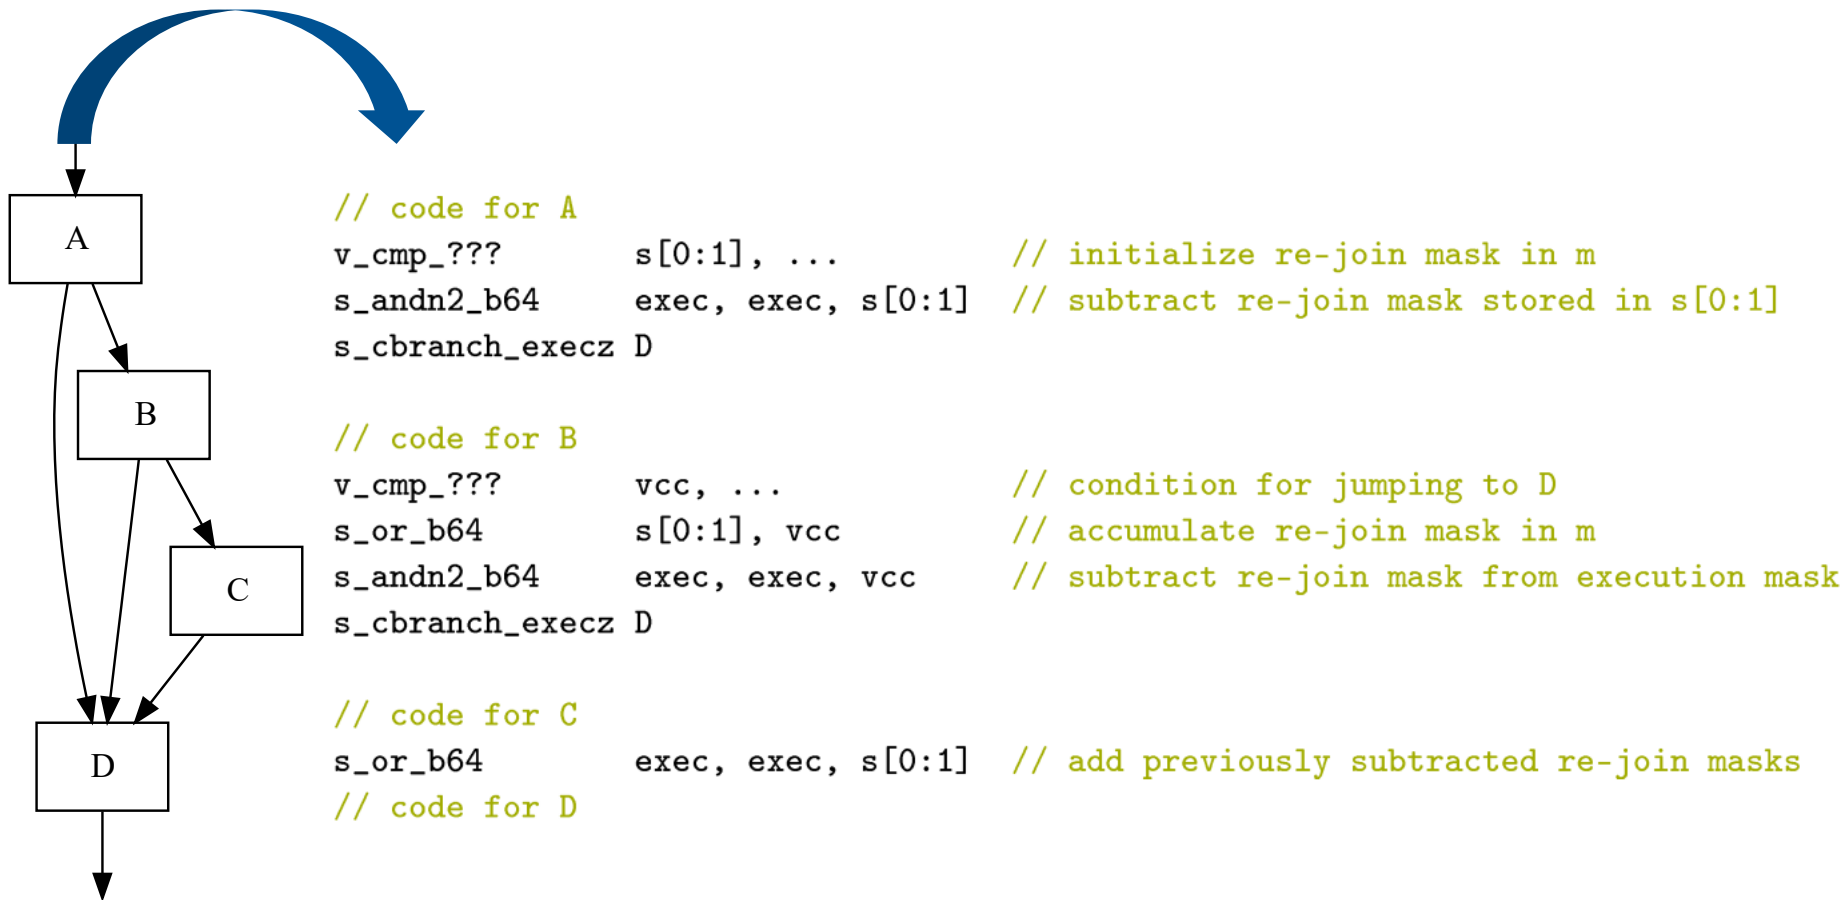
Reconverging CFGs

Definition:

- Every **non-uniform** terminator **B** (conditional branch) has exactly two successors
- One of which post-dominates **B**




Lowering to wave-level



Lowering Reconverging CFGs

For each conditional **non-uniform** node **N**:


- Virtual register **m** holds re-join mask for basic block **N** (containing diverging threads branching to the **primary** successor)
- Subtract **m** from the **exec** register and direct control flow to **secondary** successor by placing its instructions directly after basic block **N**
- Re-join mask is used to update the **exec** register at the beginning of the **primary** successor
- **m** must be correctly initialized (at entry points to divergent flow) to avoid unrelated data being merged into the execution mask



Lowering Reconverging CFGs

For each unprocessed **primary** successor **N** in the CFG:

- Decide how to initialize or modify mask **m** for **incoming edges**
- Decide how to update the **exec** mask register (binary op)
- Insert scalar branch instruction to skip **N** if **exec** is empty
 - (might skip mask initialization instructions of previous basic blocks)
- ✓ Mark **N** as processed




Lowering Reconverging CFGs

Decide how to initialize or modify mask m for incoming edges:

- P: Inspect predecessors for which N is the primary successor (B)
- P*: Inspect (direct) successors of nodes in P, capturing secondary successors and N itself (B, C, D)

For every node X in P

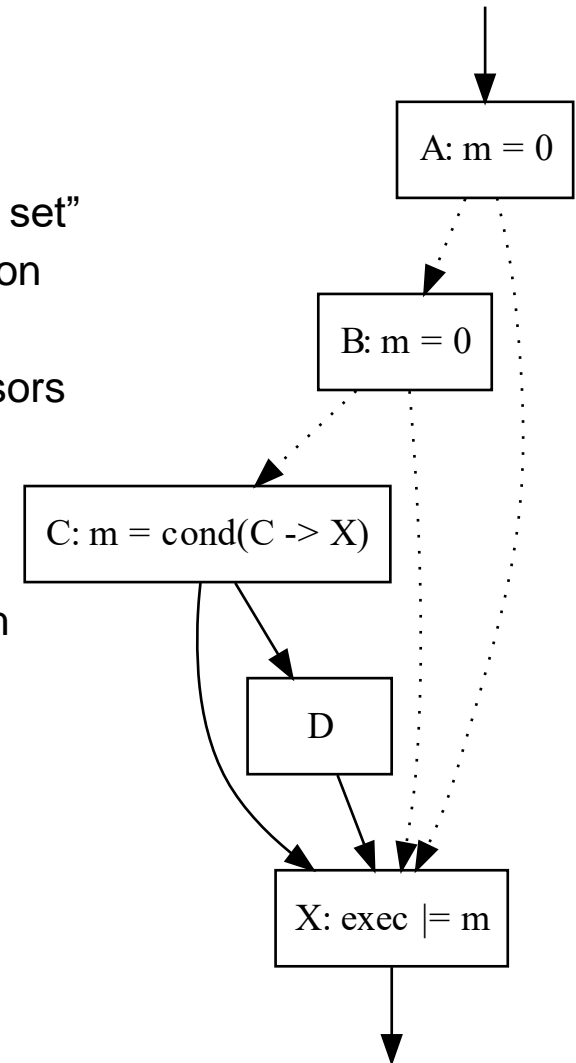
- If X has a predecessor Y not in P^* (e.g. $X = B$, $Y = A$)
 - Insert MOVE instruction at X to set m to the condition of X branching to N
- Else (nested primary)
 - Insert OR instruction at X to add condition of X branching to N to mask m
- Insert AND-NOT in X to subtract condition of X branching to N from exec
- Insert scalar branch to N if exec register is empty



Lowering Reconverging CFGs

Initialize mask m for incoming **uniform**- or **backward** edges:


- Need to identify entry nodes not included in the “reachable set” to avoid uninitialized mask m corrupting thread join operation
- P^* : **successors** of nodes in P , capturing secondary successors and N itself: “reachable set” (C, D, X)
- Q : **predecessors** of “reachable set” P^* which have predecessor that are not part of the reconverging subgraph
 - Predecessors of (D, X) not part of (C, D, X) that are not covered by initialization at the **entry** node (C: A and B)
 - Insert MOVE instruction at nodes in Q , setting m to 0



Lowering Reconverging CFGs

Initialize mask m for incoming **uniform**- or **backward** edges:


- Need to identify entry nodes not included in the “reachable set” to avoid uninitialized mask m corrupting thread join operation
- Q: **predecessors** of “reachable set” P^* which have predecessor that are not part of the reconverging subgraph
- Q and the node X for which N is the **primary** successor cover all nodes entering the flow of a particular reconverging subgraph



Lowering Reconverging CFGs

Initialize mask m for incoming **uniform-** or **backward** edges:

- Split **primary** successor with backward edge so that the backward edge is a predecessor not contained in the subgraph \rightarrow part of set $Q \rightarrow$ Initialized with 0



Transforming to reconverging control flow


Approach:

- Compute basic block ordering in which to process input CFG
- Maintain *open tree* OT structure containing unprocessed “open” edges to reroute control flow towards the *exit* node by inserting new *flow* blocks

Ordering:

- Any ordering is viable as long as the exit node comes last
- Ordering is based on traversal of the input CFG
- Quality of reconverging CFG depends on input basic block ordering


OpenTree OT



Transforming to reconverging control flow

Processing nodes:


- Nodes of the **OT** have sets of open *Incoming* and *Outgoing* edges (to nodes of the **OT**, not CFG) that need to be processed
- A node can be closed if both sets are emptied by processing
- Closed nodes are removed and their child nodes moved to its parent



Transforming to reconverging control flow

Adding node **B** to OT:


- Collapse all paths leading from the root to **B** into a single path containing all its **visited** predecessors
- Attach **B** as child to the lowest predecessor
- Close all open edges from **visited** predecessors leading to **B**
- Mark **B** as visited



Transforming to reconverging control flow


Rerouting:

- Divergent nodes are called *armed* if one of the outgoing edges has already been *closed*
 - Reroute Subtree *S*:
 - Create new *flow* basic block *X*
 - For each open outgoing edge $E = (N, B)$ of nodes $N \in S$
 - Replace E by edge (N, X)
 - Add edge (X, B) if it does not already exist
 - Add basic block *X* to OT
 - Add a Phi-instruction for every successor B of *X* containing the conditions under which each predecessor N branches to the successor



Transforming to reconverging control flow

Input CFG



```

define amdgpu_ps void @inputcfg(i32 %in_A, i32 %in_C, i32 %in_D, i32 %in_E)
{
A:
    %cc_A = icmp eq i32 %in_A, 0
    br i1 %cc_A, label %B, label %C

B:
    br label %D

C:
    %cc_C = icmp eq i32 %in_C, 0
    br i1 %cc_C, label %D, label %E


D:
    %cc_D = icmp eq i32 %in_D, 0
    br i1 %cc_D, label %D, label %F

E:
    %cc_E = icmp eq i32 %in_E, 0
    br i1 %cc_E, label %C, label %F


F:
    ret void
}
  
```

Transforming to reconverging control flow


Input CFG




OT: Added A



OT: Added B



OT: Added D







Transforming to reconverging control flow

Input **CFG**

OT: Added D

CFG: Reroute Condition 1

OT: Added FLOW0







Transforming to reconverging control flow

Input **CFG**

OT: Added C


OT: Added FLOW1

CFG: Rerouting Condition 2




Transforming to reconverging control flow

Input CFG




OT: Added E



OT: Added F

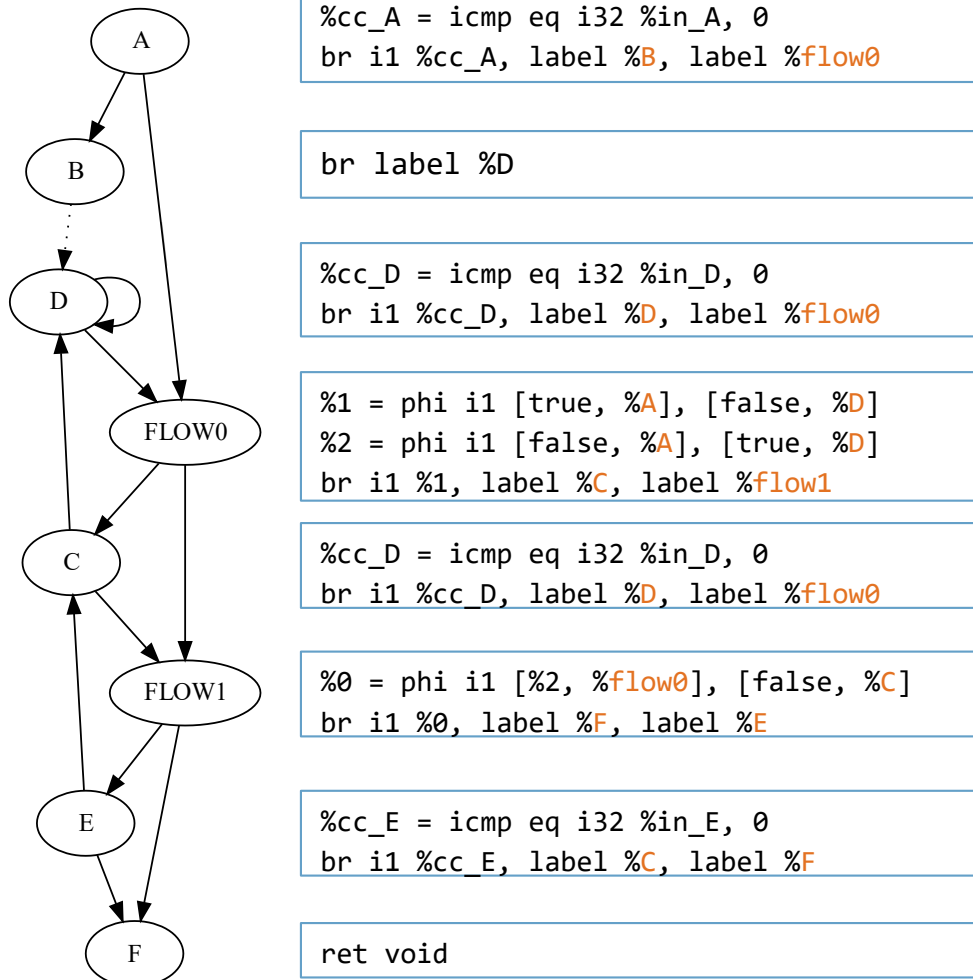


CFG: After F



Transforming to reconverging control flow

Reconverging CFG:



Reconverging Control-Flow Graphs

Contributions:

- New SPMD vectorization approach
- Simple and concise definition of Reconvergence for CFGs (weaker than structuredness)
- Proof-of-Concept lowering algorithm and CFG transformation

Properties:

- Support for unstructured and irreducible input CFGs
- Preserve uniform control flow
- Retain CFGs that are already reconverging
- Insert fewer new basic blocks than structurization requires (StructurizeCFG)