


Implementing SPMD control flow in LLVM using reconverging CFGs

Fabian Wahlster

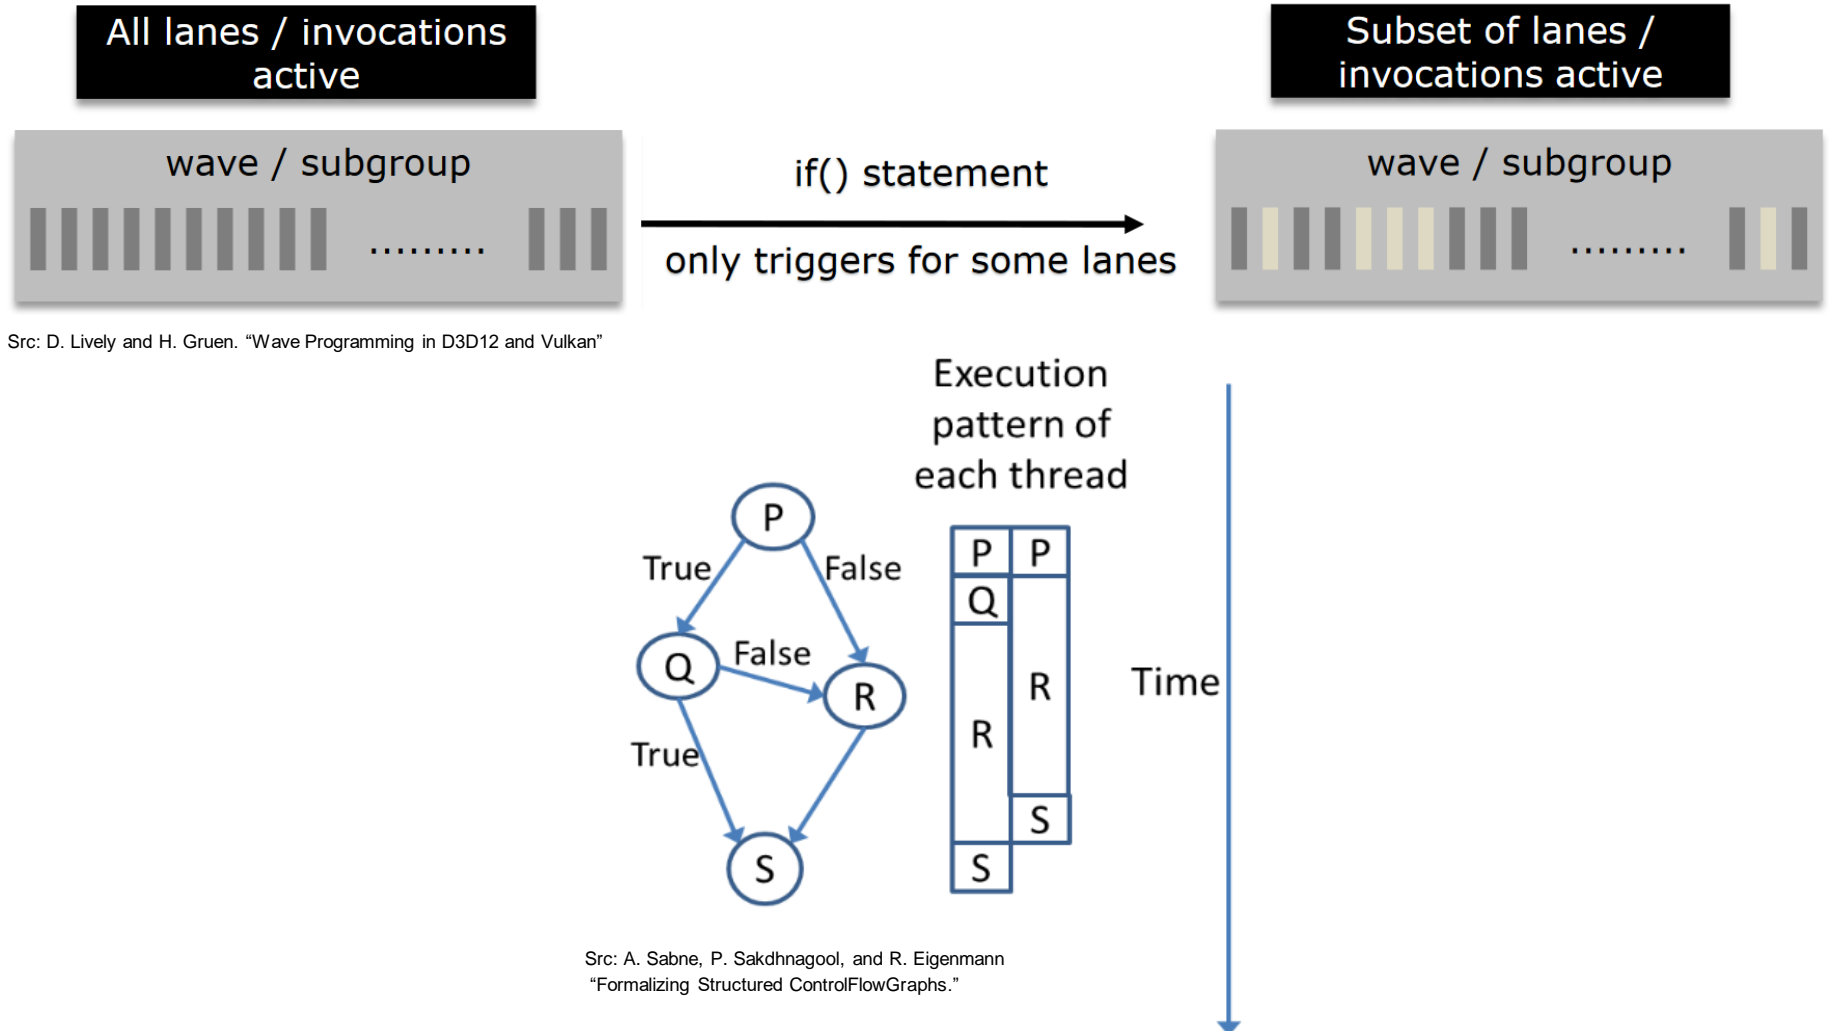
Technische Universität München – UX3D

Nicolai Hähnle

AMD



Divergence on wide SIMD





Converting thread-level code to wave-level ISA

```
float fn0(float a,float b)
{
    if(a>b)
        return ( (a-b) *a) ;
    else
        return ( (b-a) *b) ;
}
```



//Registers r0 contains "a", r1 contains "b"
//Value is returned in r2



```

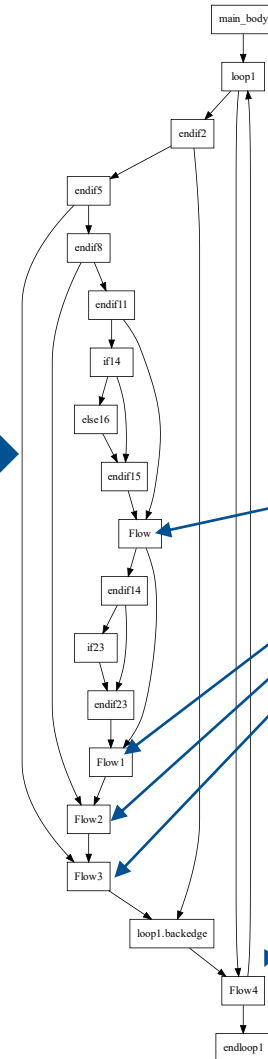

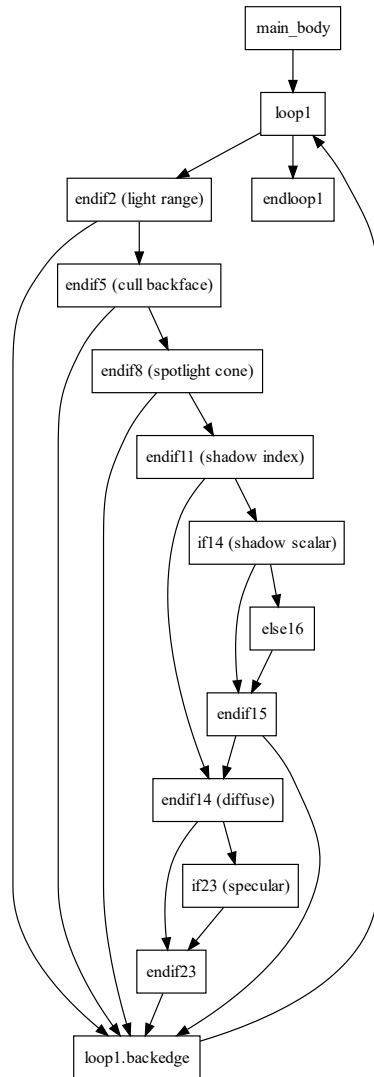
v_cmp_gt_f32    r0,r1        //a > b, establish VCC
s_mov_b64       s0,exec      //Save current exec mask
s_and_b64       exec,vcc,exec //Do "if"
s_cbranch_vccz  label0       //Branch if all lanes fail
v_sub_f32       r2,r0,r1     //result = a - b
v_mul_f32       r2,r2,r0     //result=result * a

label0:
s_andn2_b64     exec,s0,exec  //Do "else"(s0 & !exec)
s_cbranch_execz label1       //Branch if all lanes fail
v_sub_f32       r2,r1,r0     //result = b - a
v_mul_f32       r2,r2,r1     //result = result * b

label1:
s_mov_b64       exec,s0      //Restore exec mask
  
```

M. Mantor and M. Houston: AMD Graphic Core Next Architecture, Fusion 11 Summit presentation

Structurization in LLVM

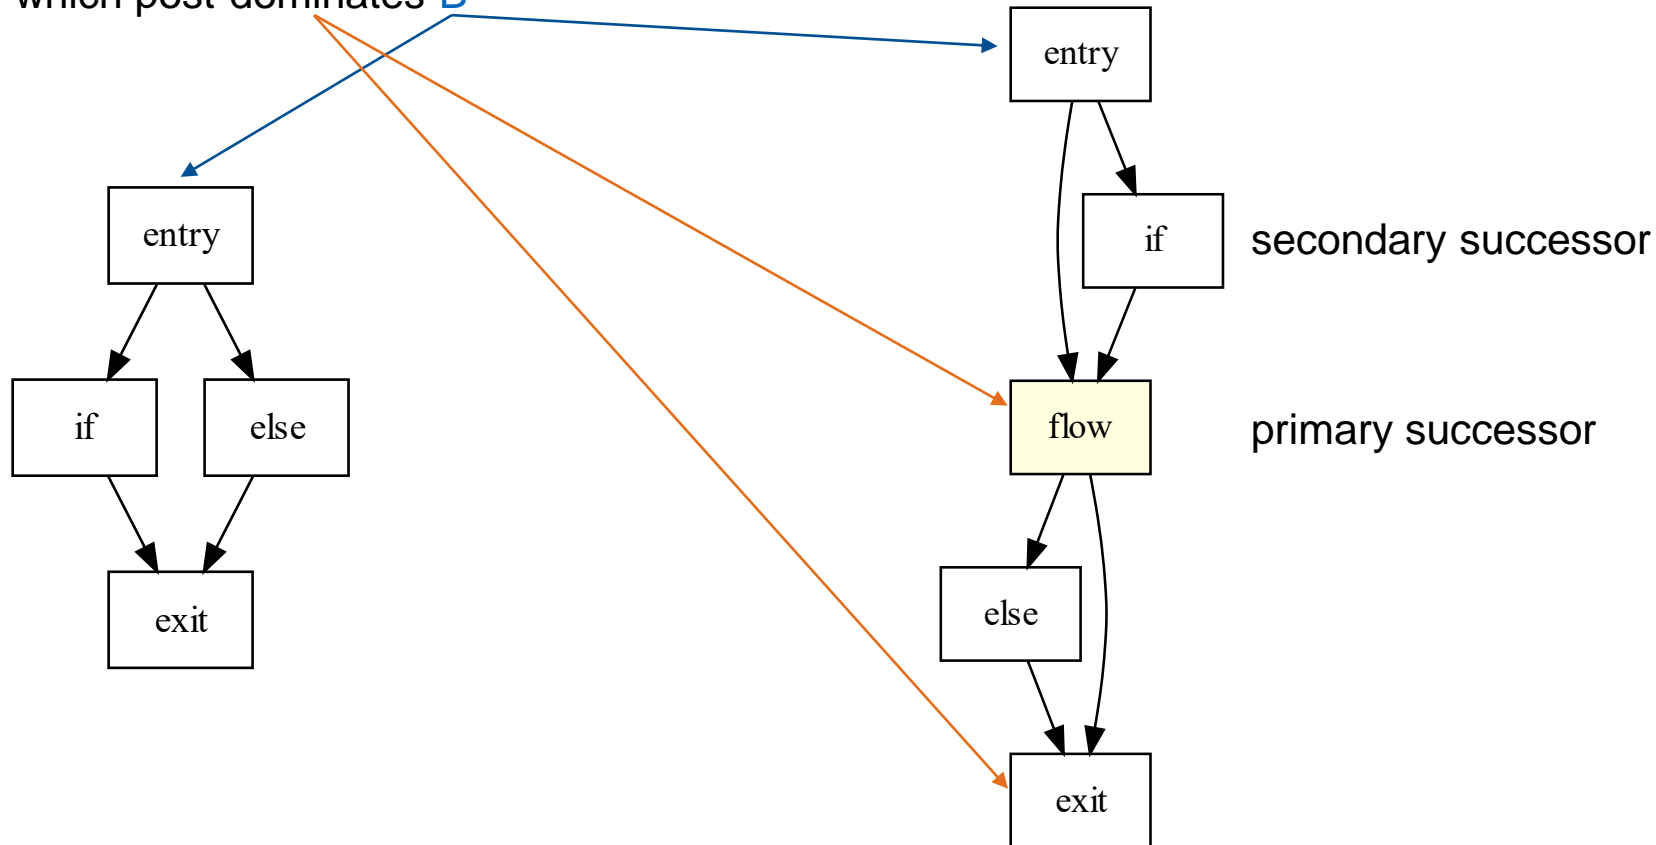


Unnecessary
flow blocks

Reconverging CFGs

Definition:


- Every **non-uniform** terminator **B** (conditional branch) has exactly two successors
- One of which post-dominates **B**



Lowering Reconverging CFGs

For each conditional **non-uniform** node **N**:

- Virtual register **m** holds re-join mask for basic block **N**
- Subtract **m** from the **exec** register to direct control flow to **secondary** successor
- Add **m** the **exec** register at the beginning of the **primary** successor to re-join divergent threads
- **m** must be correctly initialized to avoid unrelated data being merged into the execution mask



Transforming to reconverging control flow


Approach:

- Maintain *open tree* OT structure containing unprocessed *open* edges to reroute control flow towards the *exit* node by inserting new *flow* blocks

Ordering:

- Compute basic block ordering in which to process input CFG
- Ordering is based on traversal of the input CFG
- Any ordering is viable as long as the exit node comes last
- Quality of reconverging CFG depends on the input ordering

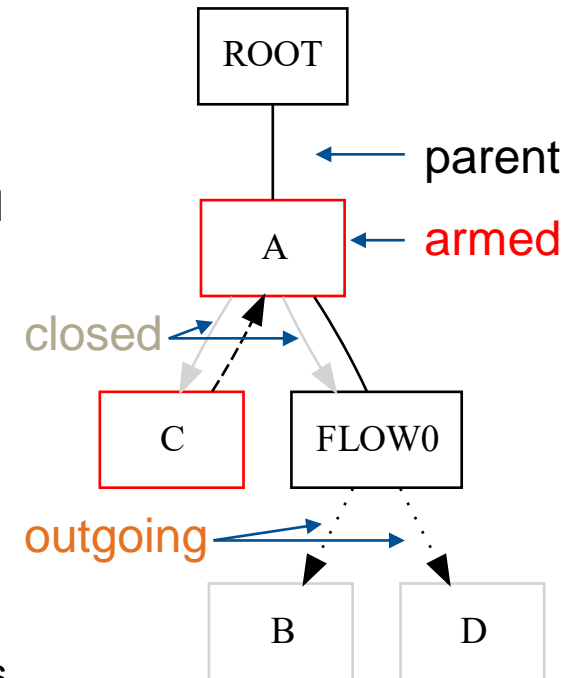
OpenTree OT



Open Tree Structure


Processing nodes:

- Nodes of the **OT** have sets of open *Incoming* and *Outgoing* edges that need to be processed
- An outgoing edge (A, B) is **closed** if A has already been visited when B is being processed
- A node can be **closed** if both sets are emptied by processing
- **Closed** nodes are removed from the **OT** and their child nodes moved to its parent
- **Divergent** nodes are called *armed* if one of the outgoing edges has already been **closed**



Transforming to reconverging control flow

Input **CFG**




$A \rightarrow C \rightarrow B \rightarrow D$


Initialize OT



Added A




Processing C...




Transforming to reconverging control flow

Input **CFG**




$A \rightarrow C \rightarrow B \rightarrow D$


Processing C...



Adding FLOW




Output **CFG**



Transforming to reconverging control flow

Reconverging CFG



A:

```
%cc_A = icmp eq i32 %in_A, 0
br i1 %cc_A, label %FLOW0, label %C
```

B:

```
br label %D
```

FLOW0:


```
%0 = phi i1 [ true, %A ], [ false, %C ]
br i1 %0, label %B, label %D
```

C:

```
%cc_C = icmp eq i32 %in_C, 0
br i1 %cc_C, label %A, label %FLOW0
```


D:

```
ret void
```





Input Ordering Exit Condition

Input CFG

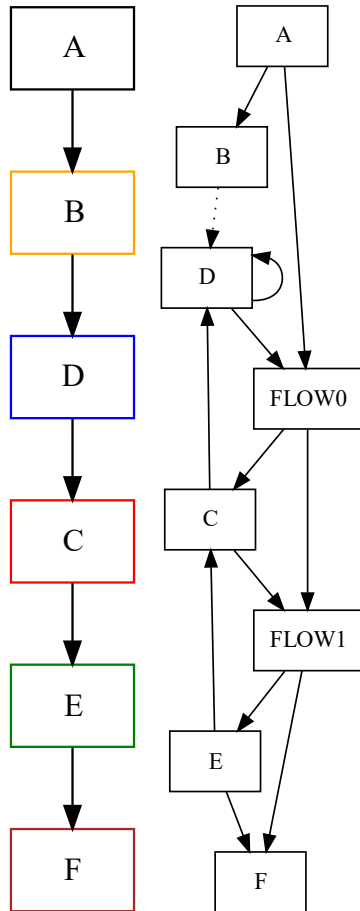


Depth First:

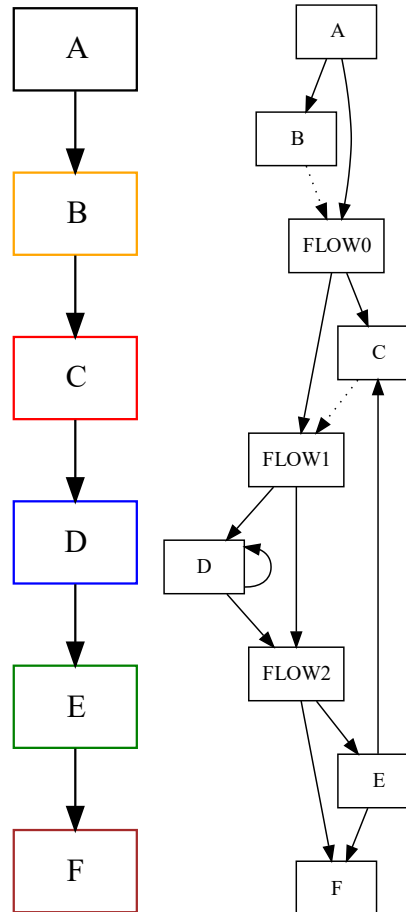


Input Ordering Comparison

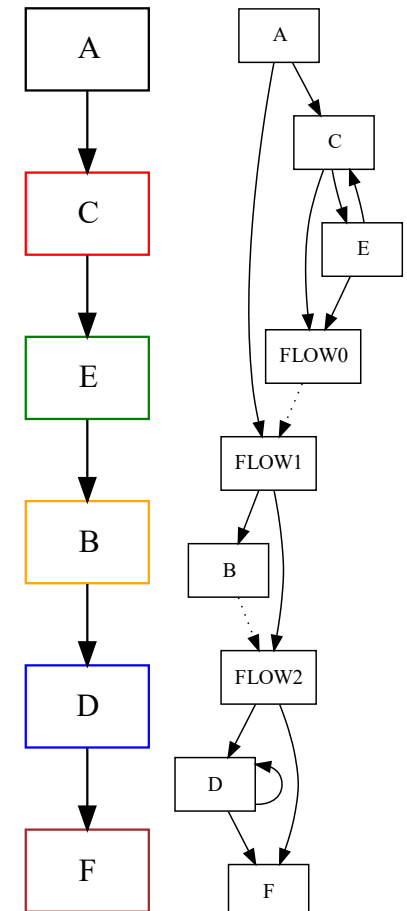
Depth First:



Breadth First:



RPOT:



Reconverging Control-Flow Graphs

Contributions:

- New SPMD vectorization approach
- Simple and concise definition of Reconvergence for CFGs (weaker than structuredness)
- Proof-of-Concept lowering algorithm and CFG transformation

Properties:

- Support for unstructured and irreducible input CFGs
- Preserves uniform control flow
- Retains CFGs that are already reconverging
- Insert fewer new basic blocks than structurization (StructurizeCFG)